

LMPool

Locality-Aware Routing and NVLink KV-Cache Transfer for Multi-GPU LLM Serving

Jialiang Li

The Hong Kong University of Science and Technology

July 28th 2026

<https://github.com/0x0addc001/LMPool>

The Placement Mismatch

What prefix caching assumes

Reuse is cheap when a request executes where its KV prefix already resides.

What data parallelism does

Each GPU owns an independent allocator and physical KV cache, while ingress assigns requests globally.

- ▶ Round robin may send a matching request to the wrong GPU.
- ▶ Strict owner routing can overload one replica.
- ▶ Unconditional KV movement can cost more than recomputation.

Design question

When should the system route to existing KV, and when should KV move to demand?

Two Principles

Routing Principle: Cache Locality

Send a request to a GPU that already owns the longest useful contiguous prefix when the saved prefill work exceeds queueing and capacity pressure.

Transfer Principle: Cache Fluidity

Copy or move complete KV block chains over a direct NVLink pair when cache placement no longer matches future demand and expected reuse repays the transaction.

Route before transfer


Transfer only when profitable

Avoid recomputation without turning locality into a load hotspot

System Architecture

LMPool: KV-Aware Multi-GPU Serving

Route for cache locality; transfer only for profitable cache fluidity



A Request Through LMPool

1. **Ingress:** LLMEngine tokenizes the prompt, builds a Sequence, and computes chained hashes for complete prompt blocks.
2. **Global route:** Global Scheduler reads versioned prefix, load, and capacity metadata and reserves the selected rank.
3. **Local admission:** The target Local Scheduler reuses ready local blocks and asks Local Block Manager to allocate missing blocks.
4. **Optional foreground transfer:** A real block shortage may trigger a cost-gated pair-local copy or move plan.
5. **Execution:** Model Runner performs prefill, writes missing KV, and then iterates decode.
6. **Publication:** The worker reports tokens and a new versioned block/load snapshot; only ready blocks become globally visible.
7. **Completion:** References are released. Complete prefixes remain cached at ref-count zero until reuse or dependency-safe reclamation.

For candidate GPU g :

$$n_g = \max(0, R - h_g),$$

$$M_g = \min(N, n_g S),$$

$$A_g = \max(0, n_g - F_g),$$

$$C_{\text{route}}(g) = Q_g + w_p M_g + w_e A_g S.$$

Decision

Select the feasible rank with minimum projected token-equivalent work.

- ▶ R : required prompt blocks; N : prompt tokens
- ▶ S : tokens per block; h_g : contiguous ready blocks
- ▶ n_g : missing blocks; M_g : missing prefill tokens
- ▶ F_g : currently free blocks
- ▶ A_g : missing blocks beyond F_g , hence reclaim pressure
- ▶ Q_g : queued, pending, and running token-equivalent work
- ▶ w_p, w_e : prefill and reclaim cost weights

Bounded spill


Bypass an overloaded owner only when additional recomputation remains bounded.

Transactional KV Transfer

KV Block Lifecycle


Local metadata controls visibility; ModelRunner owns the physical K/V tensor.

Local allocation and prefix reuse



Transactional cross-GPU transfer

entry: ready source + reserved target




Only published ready blocks are routable; reserved and received-but-hidden blocks remain invisible.

Prepare → Execute → Publish → Finalize; abort preserves the source and releases the destination reservation.

Cost-Gated Transfer Admission

Transfer Cost and Benefit Model

Admit movement only when calibrated transfer cost is lower than saved prefll work.



Static cost uses the pair-specific P95 profile plus a calibrated residual. Fixed latency is only a zero-default fallback; completed plans update EWMA residuals.

- ▶ 6 RTX 3090 GPUs, 3 direct NV4 pairs
- ▶ Qwen3-0.6B and Qwen3-1.7B, BF16
- ▶ Same per-worker KV budget
- ▶ 5 trials, fixed seed, EOS disabled
- ▶ Outputs: routing 64, memory skew 16, load skew 8 tokens
- ▶ Controlled internal ablations


Workload	Purpose
Routing	1x/3x/5x shared-prefix locality
Load skew	Background copy plus placement leases
Memory skew	Move-style capacity release

Scope

Synthetic mechanism study; not a production-trace or production-engine comparison.


Routing: Longer Prefixes Create Stable Gains

KV-Aware Routing Under Increasing Shared-Prefix Length



At 5x, routing improves throughput by **11.8% / 21.3%** and reduces TTFT by **18.6% / 25.2%** for 0.6B / 1.7B.

NVLink Data-Path Calibration




Interpretation

Every destination passed byte validation. The 1–64 block curves calibrate the pair-specific cost model; they do not alone prove serving benefit.

Transfer: Mechanism Evidence and Boundary

Skew Workloads: Throughput and First-Token Latency



Load skew: three background placements copy 87 blocks. Placement leases serve 95 later requests. TTFT falls sharply, but a higher TPOT can offset mean E2E.

Memory skew: 0.6B verifies 35 released source blocks. It remains close to routing-only; the 1.7B run does not satisfy the full offload verification predicate.

Limitations and Future Work

- ▶ Synthetic workloads and deterministic replay order; no production trace evaluation
- ▶ Internal ablations rather than production-engine comparisons
- ▶ Six RTX 3090 GPUs, two sub-2B models with identical KV geometry
- ▶ Fixed 256-token block size; no block-size sensitivity study
- ▶ Pair-specific P95 profiles and online residuals; no held-out cost-prediction error or policy ablation
- ▶ Single control process and launcher; no consensus or launcher failover

Next evaluation

Online forecast on public multi-turn traces; production baseline; 1/2/4/6-GPU and block-size scaling; cost-model prediction and admission ablations.

Takeaways

1. **Routing solves cache locality:** reuse KV where it already resides without creating a load hotspot.
2. **NVLink transfer provides cache fluidity:** move complete block chains only when reuse repays the full transaction.
3. **Transfer mechanisms are verified:** background copy, placement leases, and safe source-block release execute under their respective stressors.
4. **The boundary is part of the result:** a completed transfer does not automatically improve mean E2E or throughput.

Route when possible. Transfer when profitable.

Source: <https://github.com/0x0addc001/LMPool>

Thank You for Listening!

Questions?

Acknowledgments

Thank Prof. Kai Chen and Yijun Sun for their guidance and feedback.

Jialiang Li The Hong Kong University of Science and Technology

<https://github.com/0x0addc001/LMPool>